

Github Link: <https://github.com/siva123825/Customer-churn-prediction-using-machine-learning.git>

Project Title: Predicting Customer Churn Using Machine Learning to Uncover Hidden Patterns

1. Problem Statement:

- **Churners:** Customers who are likely to discontinue their relationship with the company.
- **Non-churners:** Customers who are likely to remain with the company.

Importance of the Problem

Addressing customer churn is vital for several reasons:

- **Revenue Retention:** Retaining existing customers is often more cost-effective than acquiring new ones.
- **Customer Lifetime Value:** Long-term customers contribute more to the company's profitability.
- **Competitive Advantage:** Understanding churn can provide insights into customer satisfaction and areas for improvement.

2. Project Objectives:

- **Data Collection and Preprocessing**
 - **Objective:** Gather comprehensive customer data, including demographics, usage patterns, transaction history, and customer service interactions.
 - **Actions:** Clean the data by handling missing values, outliers, and inconsistencies.
 - **Outcome:** A well-prepared dataset ready for analysis and modeling.
- **Feature Engineering**
 - **Objective:** Identify and create relevant features that can enhance the predictive power of the model.
 - **Actions:** Generate new variables, such as total charges or average monthly usage, and encode categorical variables appropriately.
 - **Outcome:** A set of features that effectively represent customer behavior and characteristics. [Data AI Revolution](#)
 -

- **Model Development and Training**
 - **Objective:** Build and train multiple machine learning models to predict customer churn.
 - **Actions:** Implement algorithms like Logistic Regression, Decision Trees, Random Forest, and Gradient Boosting.

- **Outcome:** A trained model capable of making predictions based on historical data.
-
- **Model Evaluation**
 - **Objective:** Assess the performance of the developed models to ensure accuracy and reliability.
 - **Actions:** Use metrics such as accuracy, precision, recall, F1-score, and AUC-ROC curve to evaluate model performance.
 - **Outcome:** Identification of the best-performing model for deployment. [ScienceDirect+4LeewayHertz - AI Development Company+4Coditudo+4](#)
- **Model Deployment and Monitoring**
 - **Objective:** Deploy the selected model into a production environment for real-time predictions.
 - **Actions:** Integrate the model with existing business systems and monitor its performance over time.
 - **Outcome:** A functional system that provides ongoing churn predictions to inform business decisions.
-
- **Business Impact Analysis**
 - **Objective:** Evaluate the impact of churn prediction on business outcomes.
 - **Actions:** Analyze metrics such as customer retention rates, revenue growth, and customer lifetime value before and after implementing churn prediction strategies.
 - **Outcome:** Quantifiable evidence of the effectiveness of churn prediction in improving business performance. [Coditudo+1neptune.ai+1](#)

By achieving these objectives, the project aims to provide businesses with actionable insights that can lead to improved customer retention, optimized resource allocation, and enhanced overall profitability.

4. Data Description:



- ## Financial & Product Engagement

- **Balance:** The current balance in the customer's account.
- **NumOfProducts:** The total number of products or services the customer holds with the company.
- **HasCrCard:** Indicates whether the customer has a company-issued credit card.
- **IsActiveMember:** Indicates if the customer is currently an active member of the service. [Home](#)

5. Data Preprocessing:

1. Data Collection

- Gather data from various sources like CRM systems, transaction history, customer service logs, etc.
- Common features include: customer ID, demographic details, service usage, payment history, customer support interaction, and churn status (target variable).

2. Data Cleaning

- **Handle missing values:** Replace with mean/median/mode, drop rows/columns, or use imputation techniques.
- **Remove duplicates:** Check for and remove duplicate records to prevent bias.
- **Correct errors:** Fix inconsistencies like wrong data formats (e.g., date formats) or incorrect labels.

3. Feature Engineering

- **Create new features:** e.g., tenure groups, average usage per month, or interaction frequency.
- **Transform variables:** e.g., log transformation for skewed numerical data.

Encode temporal variables: Create features from dates such as “days since last interaction”. **Encoding Categorical Variables**

- **Label encoding:** Converts categorical text data into numerical form (useful for ordinal variables).
- **One-hot encoding:** Creates binary columns for each category (good for nominal variables).
- **Target encoding:** Replaces category values with the mean of the target variable for that category.

5. Feature Scaling

- **Standardization** (Z-score normalization): For algorithms like SVM, KNN, or logistic regression.
- **Normalization** (Min-Max scaling): Especially useful when data is not normally distributed.

6. Handling Imbalanced Data

- **Resampling techniques:**
 - **Oversampling:** e.g., SMOTE (Synthetic Minority Over-sampling Technique)
 - **Undersampling:** Reducing the majority class.
- **Use of class weights:** For models that support it (e.g., logistic regression, random forest).

7. Train-Test Split

- Divide the data into training and test (or validation) sets to evaluate model performance fairly.
- Common ratios: 70:30, 80:20, or use **K-Fold Cross-Validation** for robust evaluation.

8. Dimensionality Reduction (Optional)

- Use **PCA** (Principal Component Analysis) or **feature selection techniques** to reduce noise and improve performance.

6. Exploratory Data Analysis (EDA):

1. Understand the Dataset

- **Objective:** Predict whether a customer will churn (i.e., stop using a service).
- **Target Variable:** Usually a binary column like `Churn` (Yes/No or 1/0).
- **Features:** Could include customer demographics, account information, usage patterns, service types, etc.

2. Data Summary

- Use `.info()`, `.describe()` to get an overview:
 - Data types
 - Missing values
 - Statistical summary of numerical features
 - Unique values for categorical features

3. Missing Value Analysis

- Identify features with missing data.
- Strategy: Drop, impute (mean, median, mode), or flag them with indicators.

Univariate Analysis

- **Numerical features:**

```
python
CopyEdit
df['MonthlyCharges'].hist(bins=20)
```

- **Categorical features:**

```
python
CopyEdit
df['Contract'].value_counts().plot(kind='bar')
```

6. Bivariate Analysis

- **Churn vs Categorical Variables:**

```
python
CopyEdit
pd.crosstab(df['Contract'], df['Churn'],
            normalize='index').plot(kind='bar', stacked=True)
```

- **Churn vs Numerical Variables:**

```
python
CopyEdit
sns.boxplot(x='Churn', y='MonthlyCharges', data=df)
```

7. Correlation Analysis

- For numerical variables:

```
python
CopyEdit
sns.heatmap(df.corr(), annot=True)
```

7. Feature Engineering :

Understanding the Domain

Before feature engineering, it's essential to understand:

- What defines churn in your business (e.g., account closure, inactivity, subscription cancellation).
- The nature of your product/service and customer lifecycle.

2. Data Sources

Churn prediction models typically pull data from:

- **Customer demographics**
- **Transactional data**
- **Service usage logs**
- **Customer service interactions**
- **Subscription history**

3. Types of Features for Churn Prediction

A. Demographic Features

- Age, gender, income level, location

Type of account (individual vs. corporate) Feature Transformation Techniques

- **Binning:** Group continuous variables (e.g., age into age groups)
- **Scaling/Normalization:** Especially for distance-based models (e.g., k-NN)
- **One-hot encoding:** For categorical variables
- **Target encoding:** For high-cardinality categorical features

5. Feature Selection Techniques

- **Univariate statistics** (Chi-square, ANOVA)
- **Model-based** (feature importance from Random Forest, XGBoost)
- **Recursive Feature Elimination (RFE)**
- **Correlation analysis:** To remove multicollinearity

6. Time-Based Features (For Sequential Data)

- **Time since last purchase/login**
- **Recency, Frequency, and Monetary (RFM) analysis**
- **Seasonal patterns:** Usage changes during holidays, weekends, etc.

7. Advanced Feature Engineering

- **Natural Language Processing:** Sentiment analysis from reviews or support tickets
- **Clustering:** Assigning cluster labels (e.g., usage pattern groups)
- **Embedding techniques:** For high-dimensional categorical data

8. Automation Tools

- **Featuretools** (Python library for automated feature engineering)
- **tsfresh** (for time series data)
- **DataRobot, H2O.ai** (AutoML platforms)

Summary Table:

-

8. Model Building:

Understanding the Problem

- **Goal:** Predict whether a customer will churn (i.e., stop using a service).
- **Type:** Supervised classification problem (binary classification: churn vs no churn).

2. Data Collection and Preprocessing

Typical dataset includes features like:

- **Customer demographics** (age, gender, location)
- **Service usage** (subscription type, usage minutes, data consumption)
- **Customer support interaction** (number of complaints, support tickets)
- **Payment information** (billing method, payment history)

Preprocessing Steps:

- Handle missing values
- Encode categorical variables (One-Hot, Label Encoding)
- Feature scaling (StandardScaler, MinMaxScaler)
- Address class imbalance (e.g., using SMOTE, oversampling, or class weights)

3. Exploratory Data Analysis (EDA)

- Visualize churn rates across different features
- Correlation analysis
- Identify important variables

4. Model Selection

Start with baseline models and move to more complex ones:

- **Logistic Regression** (baseline)

- **Decision Trees**
- **Random Forest**
- **Gradient Boosting (e.g., XGBoost, LightGBM)**
- **Support Vector Machines**
- **Neural Networks (if dataset is large)**

5. Model Training & Evaluation

Split data into:

- **Training Set**
- **Validation Set**
- **Test Set**

Use metrics like:

- **Accuracy**
- **Precision, Recall, F1-score**
- **ROC-AUC**
- **Confusion Matrix**

Apply **cross-validation** to reduce overfitting and assess generalizability.

6. Feature Importance & Model Interpretation

- Use model explainability tools (e.g., SHAP, LIME)
- Identify key drivers of churn
- Align insights with business actions

7. Deployment

- Save model using `joblib` or `pickle`
- Deploy with APIs (e.g., Flask, FastAPI)
- Monitor model performance over time

8. Monitoring and Updating

- Track churn prediction performance (via dashboards)
- Re-train model periodically with new data

9. Visualization of Results & Model Insights:

When working on a **Customer Churn Prediction** project using machine learning, the **visualization of results and model insights** is essential to:

- Understand how the model is performing.
- Gain business insights from the data and model.
- Communicate findings effectively to stakeholders.

Here's a breakdown of how to visualize results and model insights in this context:

1. Confusion Matrix

- Shows **true positives**, **true negatives**, **false positives**, and **false negatives**.
- Helps assess where the model is making errors in predicting churn vs. non-churn.

2. ROC Curve & AUC Score

- **ROC Curve** shows the trade-off between true positive rate and false positive rate.
- **AUC (Area Under Curve)** quantifies the overall ability of the model to discriminate between classes.

3. Precision-Recall Curve

- Especially helpful in **imbalanced datasets** (where churners are a minority).
- Focuses on how well the model identifies actual churners.

4. Feature Importance (Model Explainability)

- Visualize which features (e.g., tenure, usage, customer service calls) are most important for predicting churn
- SHAP (SHapley Additive exPlanations):
 - Local and global interpretability.
 - Shows individual prediction reasoning.

Churn Rate by Feature

- Bar plots of churn rate across customer segments (e.g., by contract type, monthly charges).
- Helps uncover **behavioral patterns** in churn.

6. Model Performance Metrics

- **Accuracy, Precision, Recall, F1-Score** – often summarized in a table or bar chart.
- Helpful for comparing multiple models.

7. Customer Segmentation (Clustering for Insights)

- Use t-SNE or PCA to reduce dimensions and visualize customer clusters.
- Label clusters by churn rates to discover high-risk groups.

10. Tools and Technologies Used:

Data Collection & Storage

- **Databases:** MySQL, PostgreSQL, MongoDB
- **Data Lakes / Warehouses:** Amazon S3, Google BigQuery, Snowflake
- **APIs:** REST APIs for fetching customer interaction logs, payment history, etc.

2. Data Preprocessing & Analysis

- **Python Libraries:**
 - pandas – data manipulation
 - numpy – numerical operations
 - scikit-learn – preprocessing tools (e.g., encoding, scaling)
 - missingno – missing data visualization
- **Notebook Environments:**
 - Jupyter Notebook / JupyterLab
 - Google Colab

3. Data Visualization & EDA

- matplotlib, seaborn – static visualizations
- plotly, bokeh – interactive visualizations
- Tableau, Power BI – business-friendly dashboards

4. Machine Learning Models

- **Libraries:**
 - scikit-learn – Logistic Regression, Decision Trees, Random Forests
 - XGBoost, LightGBM – advanced gradient boosting
 - TensorFlow, Keras, PyTorch – neural networks for deep learning
- **Model Selection & Tuning:**
 - GridSearchCV, RandomizedSearchCV
 - Optuna, Hyperopt for hyperparameter tuning

11. Team Members and Contributions :

K.ABIKA :[*Data cleaning, EDA*]

T.R.VIGNESH :[*Feature engineering*]

S.KIRUBANANDHAM :[*Model development*]

V.SIVA VETRIVEL:[*Documentation and reporting*]

